

EXTENDING AND COMPLETING YAUHAU

Justus Adam, supervised by Sebastian Ertel and Andres Goens

```
(defn common-friends [x y]
  (intersection (friends-of x) (friends-of y)))

(defn friends-of [x]
  (fetch (FriendsOf. x)))
```

```
(defn common-friends [x y]
  (intersection (friends-of x) (friends-of y)))
```

```
(defn friends-of [x]
  (fetch (FriendsOf. x)))
```

- Batching, Caching and Concurrency desired

```
(defn common-friends [x y]  
  (intersection (friends-of x) (friends-of y)))
```

```
(defn friends-of [x]  
  (fetch (FriendsOf. x)))
```

- Batching, Caching and Concurrency desired
- Retaining concise and straight forward code

```
(defalgo common-friends [x y]  
  (intersection (friends-of x) (friends-of y)))
```

```
(defalgo friends-of [x]  
  (fetch (mk-req (FriendsOf. x) data-source)))
```

- Batching, Caching and Concurrency desired
- Retaining concise and straight forward code
- Yauhau solves issue with minimal difference in code

Goals

Current status

- Yauhau supports a base transformation on a dataflow graph

Current status

- Yauhau supports a base transformation on a dataflow graph
- The structure of the graph does not reflect control flow

Current status

- Yauhau supports a base transformation on a dataflow graph
- The structure of the graph does not reflect control flow

Current status

- Yauhau supports a base transformation on a dataflow graph
- The structure of the graph does not reflect control flow

Tasks

- Handling control flow

Current status

- Yauhau supports a base transformation on a dataflow graph
- The structure of the graph does not reflect control flow

Tasks

- Handling control flow
- Iteration/mapping with `smap`

Current status

- Yauhau supports a base transformation on a dataflow graph
- The structure of the graph does not reflect control flow

Tasks

- Handling control flow
- Iteration/mapping with `smap`
- Conditional execution with `if`

Current status

- Yauhau supports a base transformation on a dataflow graph
- The structure of the graph does not reflect control flow

Tasks

- Handling control flow
- Iteration/mapping with `smap`
- Conditional execution with `if`
- Semantics of side effects (writes)

Table of Contents

- 1 Application
- 2 Yauhau Overview
- 3 smap rewrite
- 4 if rewrite
- 5 Generalising to Context
- 6 Side Effect Semantics
- 7 Experiments and Evaluation

Table of Contents

- 1 Application
- 2 Yauhau Overview
- 3 smap rewrite
- 4 if rewrite
- 5 Generalising to Context
- 6 Side Effect Semantics
- 7 Experiments and Evaluation

- Large scale, distributed systems (Facebook, Amazon, Twitter)
- Haskell library Haxl¹
- Clojure library Muse²
- Scala implementation Stitch by Twitter³ (closed source)
- Yauhau as an Ohua plugin⁴

¹ Simon Marlow, Louis Brandy, Jonathan Coens, and Jon Purdy. 2014. There is no fork: an abstraction for efficient, concurrent, and concise data access. ICFP '14.

² Alexey Kachayev. 2015. Reinventing Haxl: Efficient, Concurrent and Concise Data Access. <https://www.youtube.com/watch?v=T-oekV8Pwv8>

³ Jake Donham. 2014. Introducing Stitch. <https://www.youtube.com/watch?v=VVpmMfT8aYw>

⁴ Sebastian Ertel, Andrés Goens, Justus Adam and Jeronimo Castrillon. Yauhau: Concise Code and Efficient I/O Straight from Dataflow. POPL '17. In submission.

Table of Contents

- 1 Application
- 2 Yauhau Overview**
- 3 smap rewrite
- 4 if rewrite
- 5 Generalising to Context
- 6 Side Effect Semantics
- 7 Experiments and Evaluation

- Yauhou hooks into the Ohua⁵ compiler

⁵ Sebastian Ertel, Christof Fetzer, and Pascal Felber. 2015. Ohua: Implicit Dataflow Programming for Concurrent Systems. PPPJ '15

⁶ Arvind and David E. Culler. 1986. Dataflow architectures. In Annual review of computer science vol. 1

⁷ J. B. Dennis. 1974. First version of a data flow procedure language. In Programming Symposium, Proceedings Colloque sur la Programmation

- Yauhou hooks into the Ohua⁵ compiler
- Parallelisation framework based on dataflow⁶⁷

⁵ Sebastian Ertel, Christof Fetzer, and Pascal Felber. 2015. Ohua: Implicit Dataflow Programming for Concurrent Systems. PPPJ '15

⁶ Arvind and David E. Culler. 1986. Dataflow architectures. In Annual review of computer science vol. 1

⁷ J. B. Dennis. 1974. First version of a data flow procedure language. In Programming Symposium, Proceedings Colloque sur la Programmation

- Yauhau hooks into the Ohua⁵ compiler
- Parallelisation framework based on dataflow^{6,7}
- Stateful functions and algorithms (`defalgo`) as abstraction

⁵ Sebastian Ertel, Christof Fetzer, and Pascal Felber. 2015. Ohua: Implicit Dataflow Programming for Concurrent Systems. PPPJ '15

⁶ Arvind and David E. Culler. 1986. Dataflow architectures. In Annual review of computer science vol. 1

⁷ J. B. Dennis. 1974. First version of a data flow procedure language. In Programming Symposium, Proceedings Colloque sur la Programmation

```
(defalgo common-friends [x y]
  (intersection
    (friends-of x)
    (friends-of y)))

(defalgo friends-of [x]
  (fetch (mk-req (FriendsOf. x)
                data-source)))
```

- Rewrites high level algorithms


```
(defalgo common-friends [x y]
  (intersection
    (friends-of x)
    (friends-of y)))
```

```
(defalgo friends-of [x]
  (fetch (mk-req (FriendsOf. x)
    data-source)))
```


- Rewrites high level algorithms
- Operates on dataflow ir

```
(let [[x y] (algo-in)
      a (mk-req (FriendsOf. x) data-source)
      b (mk-req (FriendsOf. y) data-source)
      c (fetch a)
      d (fetch b)
      e (intersection a b)]
  e)
```

- Rewrites high level algorithms
- Operates on dataflow ir
- Traverses dataflow graph along dependencies



- Rewrites high level algorithms
- Operates on dataflow ir
- Traverses dataflow graph along dependencies
- Collect fetches, replace with accumulator



- Rewrites high level algorithms
- Operates on dataflow ir
- Traverses dataflow graph along dependencies
- Collect fetches, replace with accumulator
- Transformation is very simple and naive






Table of Contents

- 1 Application
- 2 Yauhau Overview
- 3 smap rewrite**
- 4 if rewrite
- 5 Generalising to Context
- 6 Side Effect Semantics
- 7 Experiments and Evaluation

- Split inner graph around fetch



- Split inner graph around fetch
- insert `collect` and `smap` around fetch



- Split inner graph around fetch
- insert `collect` and `smap` around fetch
- Build tree of requests






Table of Contents

- 1 Application
- 2 Yauhau Overview
- 3 smap rewrite
- 4 if rewrite**
- 5 Generalising to Context
- 6 Side Effect Semantics
- 7 Experiments and Evaluation


- Split inner graph around fetch




- Split inner graph around fetch
- Replace two fetches with one being selectively given the active request




- Split inner graph around fetch
- Replace two fetches with one being selectively given the active request
- Push result to the continuation of the active branch using the initial condition



- Problem if unequal number of fetches on branches



- Problem if unequal number of fetches on branches
- Solution: Insert extra fetches



- Problem if unequal number of fetches on branches
- Solution: Insert extra fetches
- Use NoOp (empty) requests




Table of Contents


- 1 Application
- 2 Yauhau Overview
- 3 smap rewrite
- 4 if rewrite
- 5 Generalising to Context**
- 6 Side Effect Semantics
- 7 Experiments and Evaluation

- Handle arbitrary graphs with conditionals, maps etc


- Handle arbitrary graphs with conditionals, maps etc
- These structures are often interleaved

- Handle arbitrary graphs with conditionals, maps etc
- These structures are often interleaved
- Ease handling of complex nested stacks


- Hidden graph structures unified into context




- Hidden graph structures unified into context
- Context is property of subgraphs emerging from node labels




- Hidden graph structures unified into context
- Context is property of subgraphs emerging from node labels
- Has a marker for begin and end




- Hidden graph structures unified into context
- Context is property of subgraphs emerging from node labels
- Has a marker for begin and end
- Inherited property




- Hidden graph structures unified into context
- Context is property of subgraphs emerging from node labels
- Has a marker for begin and end
- Inherited property
- Subcontexts are always fully enclosed




- Context recognition is done at compile time




- Context recognition is done at compile time
- If a context opening node is found annotate subgraph




- Context recognition is done at compile time
- If a context opening node is found annotate subgraph
- Resolve full context stack by following the parent context references



- Context recognition is done at compile time
- If a context opening node is found annotate subgraph
- Resolve full context stack by following the parent context references



- Context recognition is done at compile time
- If a context opening node is found annotate subgraph
- Resolve full context stack by following the parent context references
- Hinges on the “fully enclosing” property



- Calculate contexts for each fetch
- Contexts are unwound in order of descending nesting level
- Unwinding is interleaved


Table of Contents

- 1 Application
- 2 Yauhau Overview
- 3 smap rewrite
- 4 if rewrite
- 5 Generalising to Context
- 6 Side Effect Semantics**
- 7 Experiments and Evaluation

- Execution order depends entirely on data dependencies

```
(defalgo main []  
  (let [write-result  
        (write (WriteReq. "my-data")  
                data-source)  
        fetch-data  
        (fetch (mk-req (Payload. "constant")  
                        data-source))]  
    (compute fetch-data write-result)))
```

- Execution order depends entirely on data dependencies



- Execution order depends entirely on data dependencies
- Can be enforced using `seq` operator

```
(defalgo main []  
  (let [write-result  
        (write (WriteReq. "my-data")  
                data-source)  
        enforced  
        (seq write-result  
              (fetch (mk-req  
                      (Payload. "constant")  
                      data-source))))]  
    (compute enforced write-result)))
```


- Execution order depends entirely on data dependencies
- Can be enforced using `seq` operator
- Should be implicit from data dependencies

```
(defalgo does-read [x]
  (compute
    (fetch (mk-req (Payload. "constant")
                    data-source))
    x))
(defalgo does-write []
  (write (WriteReq. "my-data") data-source)
  ...)
(defalgo main []
  (does-read (does-write)))
```

- Exploring and evaluating different approaches

Enforcing Execution Order

- Exploring and evaluating different approaches
- Combining expected semantics and efficiency

Enforcing Execution Order

- Exploring and evaluating different approaches
- Combining expected semantics and efficiency

- Exploring and evaluating different approaches
- Combining expected semantics and efficiency

Current approach

- seq'ing fetches to algorithm boundary
- for 'unconnected' fetches following writes
- Vice versa for 'unconnected' writes following reads

- Exploring and evaluating different approaches
- Combining expected semantics and efficiency

Current approach

- seq'ing fetches to algorithm boundary
- for 'unconnected' fetches following writes
- Vice versa for 'unconnected' writes following reads

Open questions

- Always enabled?
- Scope?

Table of Contents

- 1 Application
- 2 Yauhau Overview
- 3 smap rewrite
- 4 if rewrite
- 5 Generalising to Context
- 6 Side Effect Semantics
- 7 Experiments and Evaluation**

Verifying correctness (program semantics) and performance in comparison to Haxl and Muse for:

Verifying correctness (program semantics) and performance in comparison to Haxl and Muse for:

- ① Modularised graphs (functions/algorithms)
- ② Graphs with map operations (**smap**)
- ③ Graphs with conditionals (**if**)

Verifying correctness (program semantics) and performance in comparison to Haxl and Muse for:

- ① Modularised graphs (functions/algorithms)
- ② Graphs with map operations (`smap`)
- ③ Graphs with conditionals (`if`)

This requires extensions to our random code generator to allow generation of correct code with

Verifying correctness (program semantics) and performance in comparison to Haxl and Muse for:

- ① Modularised graphs (functions/algorithms)
- ② Graphs with map operations (**smap**)
- ③ Graphs with conditionals (**if**)

This requires extensions to our random code generator to allow generation of correct code with

- ① Randomly generated functions
- ② Map operations using randomly generated functions
- ③ Conditionals, with and without forcibly prefetched branches

The end. Thank you for listening.

Slides are available at <http://static.justus.science/presentations/extending-yauhau.pdf>

- Applicative functors denote independent data fetches
- Monad bind retrieved data
- Requests are GADT's encoding result type

```
commonFriends :: Id -> Id -> GenHaxl [Id]
commonFriends x y =
    intersection <$> friendsOf x
                <*> friendsOf y



friendsOf :: Id -> GenHaxl [Id]
friendsOf = dataFetch . FriendsOf
```

- Similar syntax to Haxl
- fmap/<\$> (<\$> and <*>), flat-map (>>=)
- Uses free monad to build an AST
- Traverses AST to find fetch rounds

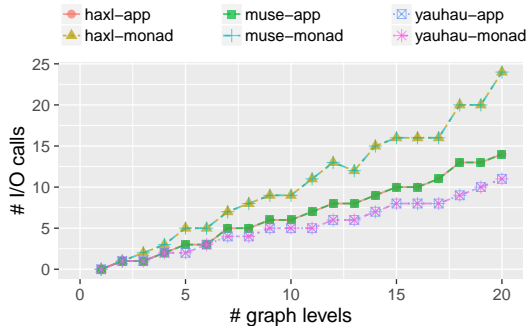

```
(defn common-friends [x y]
  (<$> intersection
    (friends-of x)
    (friends-of y)))
```

```
(defn friends-of [x]
  (FriendsOf. x))
```

```
(defrecord FriendsOf [id]
  MuseAST
  (...))
```



Left: Baseline comparison. Right: Performance with functions.



Left: Mixed latency data sources. Right: Performance across code styles.

Base transformation – Details

- Simulates program execution
- Maintains set of created data (bindings)
- At each step ‘call’ all non-fetch functions with satisfied inputs
- If no further function can be called dispatch round

```
(let [[x y] (algo-in)
      a (mk-req (FriendsOf. x) data-source)
      b (mk-req (FriendsOf. y) data-source)
      c (fetch a)
      d (fetch b)
      e (intersection a b)
      e)
```